

COMPSCI 762 Tutorial 10

Tutorial on Unsupervised Learning

Luke Chang

May 2021

The University of Auckland

Unsupervised Learning - Cluster Analysis

- Cluster analysis is an unsupervised learning task.
- The task of clustering is to partition a set of objects such that objects in the same group are more similar to each other than those in other groups.
- Evaluation metrics:
 - *Sum of Squared Error* (SSE)
 - *Sum of Squared Between*(SSB)
- Algorithms we will cover:
 - K-means
 - Hierarchical clustering
 - DBSCAN: Density-based clustering

K-means

- Partition the data into clusters
- Hyperparameter: the number of clusters, K
- Each cluster is associated with a centroid
- Each point is assigned to the cluster with the closest centroid
- Iterative method: Update centroids in each iteration, converge until the centroids don't change

Limitations:

- When clusters have:
 - Different sizes
 - Different densities
 - Non-globular shapes (hypersphere)
- When data contains outliers

K-means Pseudocode

Select K points as the initial centroids

repeat:

For each point:

 Assign the point to the closest centroid

For $i \in \{1, \dots, K\}$:

 Update the i centroid

until The centroids (or points) don't change

Evaluation Metric: Sum of Squared Error (SSE)

- No true labels are available in the unsupervised learning
- Use sum of the squared errors to evaluate the performance

$$\text{SSE} = \sum_{i=1}^K \sum_{x \in C_i} \text{dist}^2(m_i, x)$$

- SSE highly depends on **K** and the **initial centroids**
- Depends on the initial K -centroids, the K -means clusters with same the K value may not have the same SSE

Solve the Initial Centroids Problem

Problem: Stuck in the local minimal

Example

Demonstration of k-means assumptions from *sklearn*

Solution:

- Multiple runs
- Use hierarchical clustering to determine the initial centroids
- Define more than K initial centroids and then select among these initial centroids
- Postprocessing:
 - Remove empty clusters and small cluster
 - Split “loose” clusters – High SSE
 - Merge clusters when they are “close” – Low SSE
 - Apply these steps multiple times and use the pruned centroids as new initial centroids
- Improved k-means algorithms: Bisecting K-means, Mini Batch K-Means

Review Question 1: K-means

The distance matrix based on the Euclidean distance is given below:

	A1	A2	A3	A4	A5	A6	A7	A8
A1	0	$\sqrt{45}$	$\sqrt{63}$	$\sqrt{57}$	$\sqrt{41}$	$\sqrt{28}$	$\sqrt{95}$	$\sqrt{6}$
A2		0	$\sqrt{55}$	$\sqrt{49}$	$\sqrt{35}$	$\sqrt{11}$	$\sqrt{5}$	$\sqrt{25}$
A3			0	$\sqrt{11}$	$\sqrt{23}$	$\sqrt{54}$	$\sqrt{47}$	$\sqrt{65}$
A4				0	$\sqrt{2}$	$\sqrt{7}$	$\sqrt{26}$	$\sqrt{5}$
A5					0	$\sqrt{5}$	$\sqrt{21}$	$\sqrt{35}$
A6						0	$\sqrt{13}$	$\sqrt{27}$
A7							0	$\sqrt{53}$
A8								0

Suppose that the initial seeds (centers of each cluster) are **A1**, **A4** and **A7**. Run the k-means algorithm for 1 epoch only. At the end of this epoch show:

- The new clusters (i.e. the examples belonging to each cluster)
- The centers of the new clusters

Review Question 1: K-means

C1 is centred at **A1**, **C2** is centred at **A4**, and **C3** is centred at **A7**.

For each point:

- **A1 → C1**

Item	Centroid	Dist
A2	A1	$\sqrt{45}$
A2	A4	$\sqrt{49}$
A2	A7	$\sqrt{5}$

- **A2 → C3**

Item	Centroid	Dist
A3	A1	$\sqrt{63}$
A3	A4	$\sqrt{11}$
A3	A7	$\sqrt{47}$

- **A3 → C2**

- **A4 → C2**

Item	Centroid	Dist
A5	A1	$\sqrt{41}$
A5	A4	$\sqrt{2}$
A5	A7	$\sqrt{21}$

- **A5 → C2**

Item	Centroid	Dist
A6	A1	$\sqrt{28}$
A6	A4	$\sqrt{7}$
A6	A7	$\sqrt{13}$

- **A6 → C2**

- **A7 → C3**

Item	Centroid	Dist
A8	A1	$\sqrt{6}$
A8	A4	$\sqrt{5}$
A8	A7	$\sqrt{53}$

- **A8 → C2**

Review Question 1: K-means

The new clusters: $C1 = \{A1\}$, $C2 = \{A3, A4, A5, A6, A8\}$, $C3 = \{A2, A7\}$

Given $A1 = (2, 10)$, $A2 = (2, 5)$, $A3 = (8, 4)$, $A4 = (5, 8)$, $A5 = (7, 5)$, $A6 = (6, 4)$, $A7 = (1, 2)$, $A8 = (4, 9)$.

The centroids of the new clusters are:

- $C1 = (2, 10)$
- $C2 = ((8 + 5 + 7 + 6 + 4)/5, (4 + 8 + 5 + 4 + 9)/5) = (6, 6)$
- $C3 = ((2 + 1)/2, (5 + 2)/2) = (1.5, 3.5)$

Hierarchical Clustering

- Produces a set of nested clusters organized as a hierarchical tree
- Can be visualized as a dendrogram
- No assumption on the number of clusters
- Two main types:
 - Agglomerative: Start from each data point, merge the closest pair of clusters
 - Divisive: Start with one big cluster which contains all data, split at each step

Agglomerative Clustering Algorithm Pseudocode

Compute the proximity matrix

Let each point be a cluster

repeat:

 Merge the two closest clusters

 Update the proximity matrix

until only a single cluster remains

Proximity Matrix

The linkage criteria determines the metric used for the merge strategy:

- **Min (Single-linkage):** The minimum of the distances between all observations of the two sets
 - Can handle non-elliptical shapes
 - Sensitive to noise and outliers
- **Max (Complete-linkage):** The maximum distances between all observations of the two sets
 - Less susceptible to noise and outliers
 - Tends to break large clusters
 - Biased towards globular clusters
- **Group average (Average-linkage):** The average of the distances between all observations of pairs of clusters
 - Less susceptible to noise and outliers
 - Biased towards globular clusters

$$\frac{1}{|C_i| \cdot |C_j|} \sum_{i \in C_i} \sum_{j \in C_j} d(i, j)$$

where $|C_i|$ is the **cardinality** (number of elements) of the cluster C_i .

Proximity Matrix (continue)

- **Distance Between Centroids – Centroid-linkage:** Distance between the centroids of two clusters
- **Ward's minimum variance method – Ward's linkage:** Minimizes the sum of squared differences of the clusters being merged

$$\text{SSE}(C_i, C_j) - [\text{SSE}(C_i) + \text{SSE}(C_j)]$$

where $\text{SSE}(C_i, C_j)$ is the SSE of the union of the cluster i and the cluster j .

- Less susceptible to noise and outliers
- Biased towards globular clusters
- Hierarchical analogue of K-means; Can be used to initialize K-means

Example

Hierarchical clustering from *sklearn*

Review Question 2: Agglomerative Clustering

Use single-linkage (MIN) agglomerative clustering to group the data described in Exercise 1. Show the dendrogram.

	A1	A2	A3	A4	A5	A6	A7	A8
A1	0	$\sqrt{45}$	$\sqrt{63}$	$\sqrt{57}$	$\sqrt{41}$	$\sqrt{28}$	$\sqrt{95}$	$\sqrt{6}$
A2		0	$\sqrt{55}$	$\sqrt{49}$	$\sqrt{35}$	$\sqrt{11}$	$\sqrt{5}$	$\sqrt{25}$
A3			0	$\sqrt{11}$	$\sqrt{23}$	$\sqrt{54}$	$\sqrt{47}$	$\sqrt{65}$
A4				0	$\sqrt{2}$	$\sqrt{7}$	$\sqrt{26}$	$\sqrt{5}$
A5					0	$\sqrt{5}$	$\sqrt{21}$	$\sqrt{35}$
A6						0	$\sqrt{13}$	$\sqrt{27}$
A7							0	$\sqrt{53}$
A8								0

Review Question 2: Agglomerative Clustering

Use single-linkage (MIN) agglomerative clustering to group the data described in Exercise 1. Show the dendrogram.

	A1	A2	A3	A4	A5	A6	A7	A8
A1	0	$\sqrt{45}$	$\sqrt{63}$	$\sqrt{57}$	$\sqrt{41}$	$\sqrt{28}$	$\sqrt{95}$	$\sqrt{6}$
A2		0	$\sqrt{55}$	$\sqrt{49}$	$\sqrt{35}$	$\sqrt{11}$	$\sqrt{5}$	$\sqrt{25}$
A3			0	$\sqrt{11}$	$\sqrt{23}$	$\sqrt{54}$	$\sqrt{47}$	$\sqrt{65}$
A4				0	$\sqrt{2}$	$\sqrt{7}$	$\sqrt{26}$	$\sqrt{5}$
A5					0	$\sqrt{5}$	$\sqrt{21}$	$\sqrt{35}$
A6						0	$\sqrt{13}$	$\sqrt{27}$
A7							0	$\sqrt{53}$
A8								0

Level	# Clusters	Clusters
0	8	$\{A1\}, \{A2\}, \{A3\}, \{A4\}, \{A5\}, \{A6\}, \{A7\}, \{A8\}$
1	7	$\{A1\}, \{A2\}, \{A3\}, \{A4, A5\}, \{A6\}, \{A7\}, \{A8\}$

Review Question 2: Agglomerative Clustering

Use single-linkage (MIN) agglomerative clustering to group the data described in Exercise 1. Show the dendrogram.

	A1	A2	A3	A4	A5	A6	A7	A8
A1	0	$\sqrt{45}$	$\sqrt{63}$	$\sqrt{57}$	$\sqrt{41}$	$\sqrt{28}$	$\sqrt{95}$	$\sqrt{6}$
A2		0	$\sqrt{55}$	$\sqrt{49}$	$\sqrt{35}$	$\sqrt{11}$	$\sqrt{5}$	$\sqrt{25}$
A3			0	$\sqrt{11}$	$\sqrt{23}$	$\sqrt{54}$	$\sqrt{47}$	$\sqrt{65}$
A4				0	$\sqrt{2}$	$\sqrt{7}$	$\sqrt{26}$	$\sqrt{5}$
A5					0	$\sqrt{5}$	$\sqrt{21}$	$\sqrt{35}$
A6						0	$\sqrt{13}$	$\sqrt{27}$
A7							0	$\sqrt{53}$
A8								0

Level	# Clusters	Clusters
0	8	$\{A1\}, \{A2\}, \{A3\}, \{A4\}, \{A5\}, \{A6\}, \{A7\}, \{A8\}$
1	7	$\{A1\}, \{A2\}, \{A3\}, \{A4, A5\}, \{A6\}, \{A7\}, \{A8\}$
2	4	$\{A1\}, \{A2, A7\}, \{A3\}, \{A4, A5, A6, A8\}$

Review Question 2: Agglomerative Clustering

Use single-linkage (MIN) agglomerative clustering to group the data described in Exercise 1. Show the dendrogram.

	A1	A2	A3	A4	A5	A6	A7	A8
A1	0	$\sqrt{45}$	$\sqrt{63}$	$\sqrt{57}$	$\sqrt{41}$	$\sqrt{28}$	$\sqrt{95}$	$\sqrt{6}$
A2		0	$\sqrt{55}$	$\sqrt{49}$	$\sqrt{35}$	$\sqrt{11}$	$\sqrt{5}$	$\sqrt{25}$
A3			0	$\sqrt{11}$	$\sqrt{23}$	$\sqrt{54}$	$\sqrt{47}$	$\sqrt{65}$
A4				0	$\sqrt{2}$	$\sqrt{7}$	$\sqrt{26}$	$\sqrt{5}$
A5					0	$\sqrt{5}$	$\sqrt{21}$	$\sqrt{35}$
A6						0	$\sqrt{13}$	$\sqrt{27}$
A7							0	$\sqrt{53}$
A8								0

Level	# Clusters	Clusters
0	8	$\{A1\}, \{A2\}, \{A3\}, \{A4\}, \{A5\}, \{A6\}, \{A7\}, \{A8\}$
1	7	$\{A1\}, \{A2\}, \{A3\}, \{A4, A5\}, \{A6\}, \{A7\}, \{A8\}$
2	4	$\{A1\}, \{A2, A7\}, \{A3\}, \{A4, A5, A6, A8\}$
3	3	$\{A1, A4, A5, A6, A8\}, \{A2, A7\}, \{A3\}$

Review Question 2: Agglomerative Clustering

Use single-linkage (MIN) agglomerative clustering to group the data described in Exercise 1. Show the dendrogram.

	A1	A2	A3	A4	A5	A6	A7	A8
A1	0	$\sqrt{45}$	$\sqrt{63}$	$\sqrt{57}$	$\sqrt{41}$	$\sqrt{28}$	$\sqrt{95}$	$\sqrt{6}$
A2		0	$\sqrt{55}$	$\sqrt{49}$	$\sqrt{35}$	$\sqrt{11}$	$\sqrt{5}$	$\sqrt{25}$
A3			0	$\sqrt{11}$	$\sqrt{23}$	$\sqrt{54}$	$\sqrt{47}$	$\sqrt{65}$
A4				0	$\sqrt{2}$	$\sqrt{7}$	$\sqrt{26}$	$\sqrt{5}$
A5					0	$\sqrt{5}$	$\sqrt{21}$	$\sqrt{35}$
A6						0	$\sqrt{13}$	$\sqrt{27}$
A7							0	$\sqrt{53}$
A8								0

A4 and A6 are already in one cluster.


Level	# Clusters	Clusters
0	8	$\{A1\}, \{A2\}, \{A3\}, \{A4\}, \{A5\}, \{A6\}, \{A7\}, \{A8\}$
1	7	$\{A1\}, \{A2\}, \{A3\}, \{A4, A5\}, \{A6\}, \{A7\}, \{A8\}$
2	4	$\{A1\}, \{A2, A7\}, \{A3\}, \{A4, A5, A6, A8\}$
3	3	$\{A1, A4, A5, A6, A8\}, \{A2, A7\}, \{A3\}$
4	1	$\{A1, A2, A3, A4, A5, A6, A7, A8\}$

Review Question 2: Agglomerative Clustering

Use single-linkage (MIN) agglomerative clustering to group the data described in Exercise 1. Show the dendrogram.

Level	# Clusters	Clusters
0	8	$\{A1\}, \{A2\}, \{A3\}, \{A4\}, \{A5\}, \{A6\}, \{A7\}, \{A8\}$
1	7	$\{A1\}, \{A2\}, \{A3\}, \{A4, A5\}, \{A6\}, \{A7\}, \{A8\}$
2	4	$\{A1\}, \{A2, A7\}, \{A3\}, \{A4, A5, A6, A8\}$
3	3	$\{A1, A4, A5, A6, A8\}, \{A2, A7\}, \{A3\}$
4	1	$\{A1, A2, A3, A4, A5, A6, A7, A8\}$

Use the sequence from the second last level



DBSCAN: Density-based Clustering

- DBSCAN: Density-Based Spatial Clustering of Applications with Noise
- **Density:** The number of points within a specified radius (ϵ)
- **MinPts (min_samples):** A point is a **core point** if it has at least **MinPts** within ϵ .
- A **border point** is not a **core point**, but is in the neighborhood of a core point.
- A **noise point** is any point that is not a core point or a border point.

Limitations:

- Clusters with varied densities
- High dimensional data

DBSCAN Pseudocode

```
DBSCAN(DB, distFunc, eps, minPts) {  
  C := 0                                /* Cluster counter */  
  for each point P in database DB {  
    if label(P) ≠ undefined then continue /* Previously processed in inner loop */  
    Neighbors N := RangeQuery(DB, distFunc, P, eps) /* Find neighbors */  
    if |N| < minPts then {                /* Density check */  
      label(P) := Noise                  /* Label as Noise */  
    }  
    C := C + 1                            /* next cluster label */  
    label(P) := C                        /* Label initial point */  
    SeedSet S := N \ {P}                 /* Neighbors to expand */  
    for each point Q in S {               /* Process every seed point Q */  
      if label(Q) = Noise then label(Q) := C /* Change Noise to border point */  
      if label(Q) ≠ undefined then continue /* Previously processed (e.g., border point) */  
      label(Q) := C                      /* Label neighbor */  
      Neighbors N := RangeQuery(DB, distFunc, Q, eps) /* Find neighbors */  
      if |N| ≥ minPts then {              /* Density check (if Q is a core point) */  
        S := S ∪ N                        /* Add new neighbors to seed set */  
      }  
    }  
  }  
}
```

Example

Visualizing DBSCAN Clustering by Harris, 2015

Cluster Cohesion and Separation

Cluster Cohesion

- Measures how closely related are data points in a cluster
- *Within Cluster Sum of Squares* (WCSS) = *Sum of Squared Error* (SSE)

$$\text{SSE} = \sum_{i=1}^K \sum_{x \in C_i} \|x - m_i\|^2$$

Cluster Separation

- Measure how distinct or well-separated a cluster is from other clusters
- *Between cluster Sum of Squares* (BSS) a.k.a. *Sum of Squared Between* (SSB)

$$\text{SSB} = \sum_{i=1}^K |C_i| (m - m_i)^2$$

where $|C_i|$ is the size of cluster i , m is the grand mean, m_i is the mean for the cluster i .

Cluster Cohesion and Separation

- *Sum of Squares Total* (SST): The sum of squares between the n data points and the grand mean

$$SST = SSB + SSE$$

- SST is a constant based on the observed data points.
- The same terminologies are used in the one-way analysis of variance (**ANOVA**) test.

Sums of Squares (SS) Example

Example

Divide 1D data points $\{1, 2, 3, 6, 7\}$ into 2 clusters: $\{1, 2, 3\}$ and $\{6, 7\}$.

- $n = 5$, $|C_1| = 3$, and $|C_2| = 2$
- $m = (1 + 2 + 3 + 6 + 7)/5 = 3.8$
- $m_1 = (1 + 2 + 3)/3 = 2$
- $m_2 = (6 + 7)/2 = 6.5$

$$\text{SSB} = 3(3.8 - 2)^2 + 2(3.8 - 6.5)^2 = 24.3$$

$$\text{SSE} = (1 - 2)^2 + (2 - 2)^2 + (3 - 2)^2 + (6 - 6.5)^2 + (7 - 6.5)^2 = 2.5$$

$$\text{SST} = \text{SSB} + \text{SSE} = 24.3 + 2.5 = 26.8$$

The alternative way to compute SST:

$$\text{SST} = (1 - 3.8)^2 + (2 - 3.8)^2 + (3 - 3.8)^2 + (6 - 3.8)^2 + (7 - 3.8)^2 = 26.8$$

Silhouette Coefficient

Silhouette coefficient combines ideas of both cohesion and separation, but for individual points.

For each data point, i :

a: The mean distance between a sample and all other points in the same class

$$a(i) = \frac{1}{|C_i| - 1} \sum_{j \in C_i, j \neq i} d(i, j)$$

b: The mean distance between a sample and all other points in the **next nearest cluster**

$$b(i) = \min_{k \neq i} \frac{1}{|C_k|} \sum_{j \in C_k} d(i, j)$$

Silhouette Coefficient for the Entire Dataset

The Silhouette Coefficient for a single sample $s(i)$ is then given as:

$$s(i) = \begin{cases} 1 - a(i)/b(i) & , \text{ if } a(i) < b(i) \\ b(i)/a(i) - 1 & , \text{ if } a(i) \geq b(i) \end{cases}$$

- The Silhouette Coefficient for a set of samples is given as the mean of the Silhouette Coefficient for each sample.
- The range is in $[0, 1]$, the **larger** the better.

Example

An example from *sklearn*

Overview of Clustering Methods From *sklearn*

Example

The *sklearn* Documentation for Clustering